



Forecast Linear Augmented Projection (FLAP): A free lunch to reduce forecast error variance

Yangzhuoran Fin Yang
George Athanasopoulos
Rob J Hyndman
Anastasios Panagiotelis

Forecast Linear Augmented Projection (FLAP)

A model-independent post-forecast adjustment method that can reduce forecast error variance.

- Averaging indirect forecasts from linear combinations (components)
- Projecting forecasts of augmented series
- Free lunch: no additional data or information needed

What to expect

- Intuition with data
- Literature
- Method formulation
- Properties
- Empirical applications and simulation


Australian tourism data

- The data include tourism information on seven states and territories which can be divided into 77 regions
 - ▶ For example, Melbourne, Sydney, East Coast


Visitor nights

The total number of nights spent by Australians away from home recorded monthly

Melbourne and Sydney



Total and Region



Intuition

Observation

1. Similar patterns are shared by different series.
2. Better signal-noise ratio in the linear combination.

Intuition

Observation

1. Similar patterns are shared by different series.
2. Better signal-noise ratio in the linear combination.

One step further

Finding components that

1. are easy to forecast;
2. can capture the common signals;
3. can improve forecast of original series.

Literature

Forecast reconciliation

- Wickramasuriya, Athanasopoulos, and Hyndman (2019): Projecting forecasts to be consistent with the hierarchical structure

Forecast combination

- Combining forecasts of the target series
- Hollyman, Petropoulos, and Tipping (2021): Combining direct and indirect forecasts
- Petropoulos and Spiliotis (2021): Combining forecasts of selections and transformations of the target series (“wisdom of data”)

Literature


Bagging

- Bergmeir, Hyndman, and Benítez (2016): Bagging ETS models to forecast
- Petropoulos, Hyndman, and Bergmeir (2018): The benefits of bagging originate from the model uncertainty


Dynamic factor model (DFM)

- Stock and Watson (2002a), Stock and Watson (2002b)
- De Stefani et al. (2019): Machine learning extension

Series $y_t \in \mathbb{R}^m$



Components $\mathbf{c}_t = \Phi \mathbf{y}_t \in \mathbb{R}^p$



$$\mathbf{z}_t = \begin{bmatrix} \mathbf{y}_t \\ \mathbf{c}_t \end{bmatrix} \quad \tilde{\mathbf{z}}_{t+h} = \mathbf{M}\hat{\mathbf{z}}_{t+h}$$

$$\tilde{\mathbf{y}}_{t+h} = \mathbf{J}\tilde{\mathbf{z}}_{t+h} = \mathbf{J}\mathbf{M}\hat{\mathbf{z}}_{t+h}$$


$$\mathbf{M} = \mathbf{I}_{m+p} - \mathbf{W}_h \mathbf{C}' (\mathbf{C} \mathbf{W}_h \mathbf{C}')^{-1} \mathbf{C}$$

$$\mathbf{J} = \mathbf{J}_{m,p} = \begin{bmatrix} \mathbf{I}_m & \mathbf{O}_{m \times p} \end{bmatrix}$$

$$\mathbf{C} = \begin{bmatrix} -\Phi & \mathbf{I}_p \end{bmatrix}$$

$$\mathbf{W}_h = \text{Var}(\mathbf{z}_{t+h} - \hat{\mathbf{z}}_{t+h})$$

Forecasts and FLAP of series



Nonnegative variance reduction

Under unbiasedness, the variance reduction is **positive semi-definite**:

$$\begin{aligned}\text{Var}(\mathbf{y}_{t+h} - \hat{\mathbf{y}}_{t+h}) - \text{Var}(\mathbf{y}_{t+h} - \tilde{\mathbf{y}}_{t+h}) \\ = \mathbf{J}\mathbf{W}_h\mathbf{C}'(\mathbf{C}\mathbf{W}_h\mathbf{C}')^{-1}\mathbf{C}\mathbf{W}_h\mathbf{J}'\end{aligned}$$

Projection matrix


Projection matrix

The matrix \mathbf{M} is a projection onto the space where the constraint $\mathbf{C}\mathbf{z}_t = \mathbf{0}$ is satisfied.


Properties

- 1 The projected forecast $\tilde{\mathbf{z}}_{t+h}$ satisfies the constraint $\mathbf{C}\tilde{\mathbf{z}}_{t+h} = \mathbf{0}$.
- 2 For \mathbf{z}_{t+h} that already satisfies the constraint, the projection does not change its value:
 $\mathbf{M}\mathbf{z}_{t+h} = \mathbf{z}_{t+h}$
- 3 If the base forecasts are unbiased, then the FLAP forecasts are also unbiased.


Geometry of FLAP




Geometry of FLAP




Geometry of FLAP




Geometry of FLAP




Geometry of FLAP




Geometry of FLAP



Geometry of FLAP



Geometry of FLAP



Example $W_h = I_{m+p}$

$$\begin{aligned}\text{Var}(\mathbf{y}_{t+h} - \hat{\mathbf{y}}_{t+h}) - \text{Var}(\mathbf{y}_{t+h} - \tilde{\mathbf{y}}_{t+h}) \\&= \mathbf{J}\mathbf{C}'(\mathbf{C}\mathbf{C}')^{-1}\mathbf{C}\mathbf{J}' \\&= \Phi'(\Phi\Phi' + \mathbf{I})^{-1}\Phi\end{aligned}$$

Let Φ consist of orthogonal unit vectors:

$$\Phi\Phi' = \mathbf{I}_p \text{ when } p \leq m$$

$$\Phi'\Phi = \mathbf{I}_m \text{ when } p = m.$$

$$\begin{aligned}\text{tr}(\text{Var}(\mathbf{y}_{t+h} - \hat{\mathbf{y}}_{t+h}) - \text{Var}(\mathbf{y}_{t+h} - \tilde{\mathbf{y}}_{t+h})) \\&= \frac{1}{2}\text{tr}(\Phi'\Phi) = \frac{1}{2}p\end{aligned}$$

Positive condition

For the first component to have a guaranteed reduction of forecast error variance, the following condition must be satisfied:

$$\phi_1 \mathbf{W}_{y,h} \neq \mathbf{w}_{c_1 y, h},$$

- ϕ_1 is the weight vector of the first component
- $\mathbf{W}_{y,h} = \text{Var}(\mathbf{y}_{t+h} - \hat{\mathbf{y}}_{t+h})$
- $\mathbf{w}_{c_1 y, h}$ is the forecast error covariance between the first component and the original series.

Monotonicity

The forecast error variance reductions, i.e. the diagonal elements of

$$\begin{aligned} \text{Var}(\mathbf{y}_{t+h} - \hat{\mathbf{y}}_{t+h}) - \text{Var}(\mathbf{y}_{t+h} - \tilde{\mathbf{y}}_{t+h}) \\ = \mathbf{J}\mathbf{W}_h\mathbf{C}'(\mathbf{C}\mathbf{W}_h\mathbf{C}')^{-1}\mathbf{C}\mathbf{W}_h\mathbf{J}' \end{aligned}$$

is non-decreasing as p increases.

Minimum variance of individual series

The projection is equivalent to the mapping

$$\tilde{\mathbf{y}}_{t+h} = \mathbf{G}\hat{\mathbf{z}}_{t+h},$$

where $\mathbf{G} = [\mathbf{g}_1 \ \mathbf{g}_2 \ \dots \ \mathbf{g}_m]' \in \mathbb{R}^{m \times (m+p)}$ is the solution to

$$\arg \min_{\mathbf{G}} \mathbf{G}\mathbf{W}_h\mathbf{G}' \quad \text{s.t. } \mathbf{G}\mathbf{S} = \mathbf{I}$$

or

$$\arg \min_{\mathbf{g}_i} \mathbf{g}_i'\mathbf{W}_h\mathbf{g}_i \quad \text{s.t. } \mathbf{g}_i'\mathbf{s}_j = \mathbf{1}(i=j),$$

$$\text{where } \mathbf{S} = \begin{bmatrix} \mathbf{I}_m \\ \Phi \end{bmatrix} = [\mathbf{s}_1 \ \dots \ \mathbf{s}_m].$$

Key results

- 1 The forecast error variance is **reduced** with FLAP
- 2 The forecast error variance **monotonically** decreases with increasing number of components
- 3 The forecast projection is **optimal** to achieve minimum forecast error variance of each series

In practice, we need to

$$\tilde{\mathbf{y}}_{t+h} = \mathbf{J}\tilde{\mathbf{z}}_{t+h} = \mathbf{J}\mathbf{M}\hat{\mathbf{z}}_{t+h}$$

$$\mathbf{M} = \mathbf{I}_{m+p} - \mathbf{W}_h\mathbf{C}'(\mathbf{C}\mathbf{W}_h\mathbf{C}')^{-1}\mathbf{C}$$

$$\mathbf{W}_h = \text{Var}(\mathbf{z}_{t+h} - \hat{\mathbf{z}}_{t+h})$$

$$\mathbf{C} = \begin{bmatrix} -\Phi & \mathbf{I}_p \end{bmatrix}$$

- Estimate \mathbf{W}_h
- Construct Φ

Estimation of \mathbf{W}_h

Shrinking variance towards their median (Opge-Rhein and Strimmer 2007) and **shrinking covariance** towards zero (Schäfer and Strimmer 2005).

The shrinkage estimator is

- Positive definite, and
- Numerically stable.

In empirical applications, we assume

$$\widehat{\mathbf{W}}_h^{shr} = \eta_h \widehat{\mathbf{W}}_1^{shr}.$$

Construction of Φ

Principal component analysis (PCA)


Finding the weights matrix so that the resulting components **maximise variance**

Simulation


Generating values from a random distribution and normalising them to unit vectors

- Normal distribution
- Uniform distribution
- Orthonormal matrix (Borchers 2023)

Tourism (ETS)




FRED-MD (DFM)



Simulation

- Data generating process (DGP): VAR(3) with $m = 70$ variables
- Sample size: $T = 400$
- Number of repeated samples: 220
- Base model: ARIMA and DFM

Simulation



R Package flap

You can install the stable version from CRAN

```
## CRAN.R-project.org/package=flap  
install.packages("flap")
```

or the development version from Github

```
## github.com/FinYang/flap  
# install.packages("remotes")  
remotes::install_github("FinYang/flap")
```

Working paper and contact

- yangzhuoranyang.com/publication/flap/



- Email: Fin.Yang@monash.edu
- yangzhuoranyang.com

References i

- Bergmeir, Christoph, Rob J Hyndman, and José M Benítez. 2016. "Bagging Exponential Smoothing Methods Using STL Decomposition and Box-Cox Transformation." *International Journal of Forecasting* 32 (April): 303–12.
<https://doi.org/10.1016/j.ijforecast.2015.07.002>.
- Borchers, Hans W. 2023. *pracma: Practical Numerical Math Functions*.
<https://doi.org/10.32614/CRAN.package.pracma>.
- De Stefani, Jacopo, Yann-Aël Le Borgne, Olivier Caelen, Dalila Hattab, and Gianluca Bontempi. 2019. "Batch and Incremental Dynamic Factor Machine Learning for Multivariate and Multi-Step-Ahead Forecasting." *International Journal of Data Science and Analytics* 7 (June): 311–29.
<https://doi.org/10.1007/s41060-018-0150-x>.
- Hollyman, Ross, Fotios Petropoulos, and Michael E Tipping. 2021. "Understanding Forecast Reconciliation." *European Journal of Operational Research* 294 (October): 149–60. <https://doi.org/10.1016/j.ejor.2021.01.017>.
- Opgen-Rhein, Rainer, and Korbinian Strimmer. 2007. "Accurate Ranking of Differentially Expressed Genes by a Distribution-Free Shrinkage Approach." *Statistical Applications in Genetics and Molecular Biology* 6 (February).
<https://doi.org/10.2202/1544-6115.1252>.

References ii

- Petropoulos, Fotios, Rob J Hyndman, and Christoph Bergmeir. 2018. "Exploring the Sources of Uncertainty: Why Does Bagging for Time Series Forecasting Work?" *European Journal of Operational Research* 268 (July): 545–54. <https://doi.org/10.1016/j.ejor.2018.01.045>.
- Petropoulos, Fotios, and Evangelos Spiliotis. 2021. "The Wisdom of the Data: Getting the Most Out of Univariate Time Series Forecasting." *Forecasting* 3 (June): 478–97. <https://doi.org/10.3390/forecast3030029>.
- Schäfer, Juliane, and Korbinian Strimmer. 2005. "A Shrinkage Approach to Large-Scale Covariance Matrix Estimation and Implications for Functional Genomics." *Statistical Applications in Genetics and Molecular Biology* 4 (November). <https://doi.org/10.2202/1544-6115.1175>.
- Stock, James H, and Mark W Watson. 2002a. "Macroeconomic Forecasting Using Diffusion Indexes." *Journal of Business & Economic Statistics* 20 (April): 147–62. <https://doi.org/10.1198/073500102317351921>.
- . 2002b. "Forecasting Using Principal Components from a Large Number of Predictors." *Journal of the American Statistical Association* 97 (December): 1167–79. <https://doi.org/10.1198/016214502388618960>.

References iii

Wickramasuriya, Shanika L, George Athanasopoulos, and Rob J Hyndman. 2019. "Optimal Forecast Reconciliation for Hierarchical and Grouped Time Series Through Trace Minimization." *Journal of the American Statistical Association* 114 (April): 804–19. <https://doi.org/10.1080/01621459.2018.1448825>.